The file `Magazine.CSV` contains data on advertising costs and characteristics of magazines. The response variable is `PageCost`, which represents the cost of a full-page color ad in the magazine. `Circ` is the circulation of the magazine (in thousands), `MedIncome` is the median income of the readers, and `%Male` is the percentage of the readers who are male. The square root of the circulation is given in `SqrtCirc`.

In this problem, we will fit a regression model to predict the mean page cost of a magazine with `Circ` of 10000, `MedIncome` of $40,000, and `%Male` of 50. We will not necessarily use all three predictors, only the ones that are useful for predicting `PageCost`.

(a) First, we will find an appropriate set of predictor variables.

   (i) Run a multiple regression of `PageCost` on the original predictor variables (`Circ`, `MedIncome` and `%Male`). Before running it, click on Graphs, and check the box for Residuals plots: Four in one. Note that the residuals versus fit plot shows structure: a generally upward-sloping pattern, with three outliers at the right dragging things down. Identify the Magazines corresponding to the three outliers (all of which have a very large circulation).

   (ii) To investigate further, generate a scatterplot of `PageCost` versus `Circ`. Note that the plot is "bunched up" at the left, and "stretched out" at the right, and also a bit curved. In what way do the points identified as outliers in (a) deviate from the pattern in the plot here?

   (iii) To try to improve the linear relationship, let's try working with the square root of Circulation (`SqrtCirc`) rather than the circulation itself. Plot `PageCost` versus `SqrtCirc`. Based on the plot, explain why it seems more appropriate to use `SqrtCirc` as an explanatory variable in a linear regression rather than `Circ`.

   (iv) Now, run a multiple regression of `PageCost` on `SqrtCirc`, `MedIncome` and `%Male`. Plot the residuals versus fitted values. Does it look better than in (i)?

(b) Next we will investigate the regression model of `PageCost` on `SqrtCirc`, `MedIncome` and `%Male`.

   (i) Based on the p-value for the Analysis of Variance F test for this model, does the regression seem to be useful for predicting `PageCost`? Does this mean that all variables are useful?

   (ii) Which coefficients in the regression are statistically significant?

   (iii) Based on the p-values for the regression coefficients, which variables seem to be useless for predicting `PageCost`?

(c) Now, we will try to simplify the model by deleting useless predictors. Re-run the regression for `PageCost`, this time with just the two significant explanatory variables you found in part (b).

   (i) Did the R^2 go down by much compared to the regression in (b)? Is the F-statistic still significant? What does this suggest about the deleted predictor variable?

   (ii) Are the coefficients of both variables statistically significant?

(d) Finally, we will use the simplified multiple regression model to predict `PageCost`.

    (i) Get a 95% confidence interval for the mean page cost of a magazine with a `SqrtCirc` of 100, and a median income of $40,000. To do this, after running the regression click on *Stat* $\Rightarrow$ *Regression* $\Rightarrow$ *Regression* $\Rightarrow$ *Predict*. Then, enter 100 in the first line under `SqrtCirc` and enter 40000 in the first line under `MedIncome`.

    (ii) Report the 95% prediction interval.

    (iii) What is the difference between the prediction interval and the confidence interval?