

# Fast Moment-Based Estimation for Hierarchical Models

Patrick O. Perry  
New York University

# Motivating Application: Recommender Systems

- Large population of Users and Items
- Goal: Recommend items to users
- Examples: online shopping ("you might also like..."), targeted advertising

# Items

movie		title	genre
1		Toy Story (1995)	Comedy
2		Jumanji (1995)	Children
3		Grumpier Old Men (1995)	Comedy
4		Waiting to Exhale (1995)	Drama
5		Father of the Bride Part II (1995)	Comedy
.		.	.
.		.	.
.		.	.
10677		Bedtime Stories (2008)	Children
10678		Manhattan Melodrama (1934)	Drama
10679		Choke (2008)	Comedy
10680		Revolutionary Road (2008)	Drama
10681		Blackadder Back & Forth (1999)	Comedy

# Users

user
1
2
3
4
5
.
.
.
.
.
69873
69874
69875
69876
69877
69878

# Ratings

	user	movie	score	time
1	36072	21	3	1995-01-09 11:46:49
2	36072	47	5	1995-01-09 11:46:49
3	36072	1058	3	1995-01-09 11:46:49
4	34294	1	4	1996-01-29 00:00:00
5	34294	10	4	1996-01-29 00:00:00
.	.	.	.	.
.	.	.	.	.
.	.	.	.	.
10000050	61718	2395	3.5	2009-01-05 04:52:12
10000051	61718	6887	3.5	2009-01-05 04:52:17
10000052	61718	2869	2	2009-01-05 04:52:22
10000053	61141	4691	2.5	2009-01-05 04:55:03
10000054	61141	9153	3	2009-01-05 05:02:16

# Two Main Approaches

**Content-based:** recommend items similar to those the user liked in the past

**Collaborative:** recommend items that similar users liked

# A Model That Does Both

Group (User)  $i = 1, \dots, M$ :

- predictors  $X_i$  ( $n_i \times p$ )
- response  $y_i$  ( $n_i$ )

Model:

$$E(y_{ij} \mid \beta_i) = x_{ij}^T \beta_i$$

$$\beta_i \sim N(\mu, \Sigma)$$

(Condliff et al. 1999; Ansari et al. 2000)

# Response

$$y_i = \begin{bmatrix} y_{i1} \\ y_{i2} \\ \vdots \\ y_{in_i} \end{bmatrix}$$

# Features

$$X_i = \begin{bmatrix} X_{i11} & X_{i12} & \cdots & X_{i1p} \\ X_{i21} & X_{i22} & \cdots & X_{i2p} \\ \vdots & \vdots & \ddots & \vdots \\ X_{in_i1} & X_{in_i2} & \cdots & X_{in_ip} \end{bmatrix}$$

```
rating
5
3
5
3
3
3
5
.
.
.
```

```
action children comedy drama movie.popularity user.liked.last
0          0          0          1          1.1          0
0          0          0          1          0.1          0
1          0          0          0          1.8          0
1          0          0          0          -0.8         1
0          0          0          1          0.3          0
0          0          1          0          0.4          0
.          .          .          .          .          .
.          .          .          .          .          .
.          .          .          .          .          .
```

## User $i$ 's expected rating for item $j$ :

$$E(y_{ij} \mid \beta_i) = x_{ij}^T \beta_i$$

Item attributes




The diagram consists of two arrows pointing towards the term  $x_{ij}^T$  in the equation above. One arrow originates from the text 'Item attributes' and points to the  $x_{ij}$  part of the term. The other arrow originates from the text 'User tastes' and points to the  $\beta_i$  part of the term.

User tastes

### Example:

*"User  $i$  likes action movies,  
movie  $j$  is an action movie;  
we should recommend movie  $j$  to user  $i$ ."*

## User $i$ 's tastes:



## Example:

*"User  $i$  likes action movies,  
users who like action tend to dislike drama;  
we should assume that user  $i$  dislikes drama."*



# Hierarchical Model

Content-based:

$$E(y_{ij} \mid \beta_i) = x_{ij}^T \beta_i$$

Collaborative:

$$\beta_i \sim N(\mu, \Sigma)$$

# Problem: Fitting at Commercial Scale

(Zhang and Koren, 2007; Agarwal, 2008;  
Naik et al., 2008; Agarwal and Chen, 2009)

# Likelihood-Based Fitting is Slow

Method	Initial Cost	Cost per Iteration	Iterations
Expectation-Maximization	$Np^2$	$Mp^3$	Hundreds
Newton-Raphson	$Np^2$	$Mp^4$	Tens
Profile Likelihood (lme4)	$Np^2$	$Mp^3$	Tens to Hundreds

(Movielens:  $M \approx 10^5$ ,  $N \approx 10^7$ ,  $p \approx 10$ )

# Popular Approach #1: Split/Combine

**Idea:** divide data between  $K$  processors, compute separate estimates, then combine

**Pro:** cuts wall clock time by a factor of  $K$  (but does not reduce total amount of computation)

(Huang and Gelman, 2005; Gebregziabher et al. 2012; Scott et al. 2013, ...)

# Popular Approach #2: Stochastic Gradient Descent (SGD)

**Idea:** maximize h-likelihood (treating random effects like parameters), use gradient-based optimization

**Pro:** often faster than maximum likelihood

**Con:** requires tuning parameters, can sometimes be inconsistent

(Lee and Nelder 2006; Dror et al. 2011, ...)

# Today's Talk: Moment-Based Estimation

**Idea:** split data into  $M$  chunks, compute group-specific effect estimates, then use moment matching for population parameters


**Pro:** non-iterative (typically faster than ML), trivially parallelizable

**Con:** loss in statistical efficiency (sometimes)

# Comparison

Method	Initial Cost	Cost per Iteration	Iterations
Expectation-Maximization	$Np^2$	$Mp^3$	Hundreds
Newton-Raphson	$Np^2$	$Mp^4$	Tens
Profile Likelihood (lme4)	$Np^2$	$Mp^3$	Tens to Hundreds
Stochastic Gradient Descent	0	$Np$	Tens
Moment-Based	$Np^2$	$Mp^3$	2

# Computation Time





# Remainder of the Talk

1. **Moment-based as fast alternative to likelihood-based estimation**
2. **Consistent, asymptotically normal**
3. **Performs well in practice**

# Moment-based estimation: A new (old) estimation method, dramatically faster

History: Cochran (1937), Yates and Cochran (1938),  
Cochran (1954), Swamy (1970), Carter and Yang (1986),  
Cox and Solomon (2002)

# Intuition for Moment-Based Estimation

1. Compute group-specific coefficient estimates
2. Estimate population parameters by matching coefficient moments.

# Intuition for Moment-Based Estimation (Details)

## 1. Model:

$$y_i = X_i \beta_i + \varepsilon_i \quad \beta_i \sim N(\mu, \Sigma), \quad \varepsilon_i \sim N(0, \sigma^2 I)$$

## 2. Group-specific coefficient estimates:

$$b_i = (X_i^\top X_i)^{-1} X_i^\top y_i$$

## 3. Moments:

$$E(b_i) = \mu \quad \text{Cov}(b_i) = \Sigma + \sigma^2 (X_i^\top X_i)^{-1}$$

# Problem: Rank-Degenerate $X$

**Group-specific coefficient estimates:**

$$b_i = (X_i^T X_i)^\dagger X_i^T y_i$$

**Biased estimate:**

$$E(b_i \mid \beta_i) \neq \beta_i$$

# Solution for Rank-Degenerate $X$

1. Group-specific coefficient estimates:

$$b_i = (X_i^T X_i)^\dagger X_i^T y_i \quad (\text{biased in nullspace of } X_i)$$

2. Choose weight matrices:

$$W_i \quad (\text{same nullspace as } X_i)$$

3. Moments:

$$E(W_i b_i) = W_i \mu \quad \text{Cov}(W_i b_i) = W_i \{ \Sigma + \sigma^2 (X_i^T X_i)^\dagger \} W_i^T$$

# Moment Matching for Mean

$$\hat{\mu} = \Omega_1^{-1} \sum_{i=1}^M W_i b_i$$

$$\Omega_1 = \sum_{i=1}^M W_i$$

# Moment Matching for Covariance

$$\hat{A} = \sum_{i=1}^M W_i (b_i - \hat{\mu})(b_i - \hat{\mu})^T W_i$$

$$\text{vec}(\hat{\Sigma}) = \Omega_2^{-1} \text{vec}(\hat{A}) - \text{Bias}$$


$$\Omega_2 = \sum_{i=1}^M W_i \otimes W_i$$



# Optimal Weights

$$W_i = U_i [U_i^T \{ \Sigma + \sigma^2 (X_i^T X_i)^\dagger U_i \}^{-1} U_i^T]$$

Orthonormal columns,  
same span as  $X_i^T$



Unknown parameters



# Computational Complexity

- Compute  $M$  group-specific estimates:  $O(Np^2)$
- Match weighted moments:  $O(Mp^3)$

# Theoretical Properties

# **Theorem 1: Moment-Based Estimates are Consistent**

# Theorem 1: Details

**Statement:**

$$\hat{\mu} = \mu + O_P(\|\Omega_1^{-1}\|^{1/2})$$

$$\hat{\Sigma} = \Sigma + O_P(\|\Omega_2^{-1}\|^{1/2})$$

**Main Assumption:**

**group-specific estimates have finite fourth moments**

**Proof:**

**Linear Algebra + Markov's Inequality**

**Theorem 2: Two-Step Moment Based  
Estimates are Asymptotically  
Relatively Efficient**

# Theorem 2: Details

**Statement:**

$$\hat{\mu}_{\text{two-step}} = \hat{\mu}_{\text{optimal}} + o_P(\|\Omega_1^{-1}\|^{1/2})$$

**Assumptions:**

(same as Theorem 1)

**Proof:**

Taylor expansion / Matrix inverse perturbation

**Theorem 3: Two-Step Moment Based  
Estimates are Asymptotically Normal**



# Theorem 3: Details

**Statement:**

$$\Omega_1^{1/2}(\hat{\mu} - \mu) \implies \mathcal{N}(0, I)$$

**Assumptions:**

(same as Theorem 1 + large  $M$ )

**Proof:**


Central Limit Theorem

# Theory: Recap


- Estimates are consistent
- Two-step estimates are asymptotically relatively efficient
- Estimates are asymptotically normal

# Performance in Hierarchical Linear Model Simulations


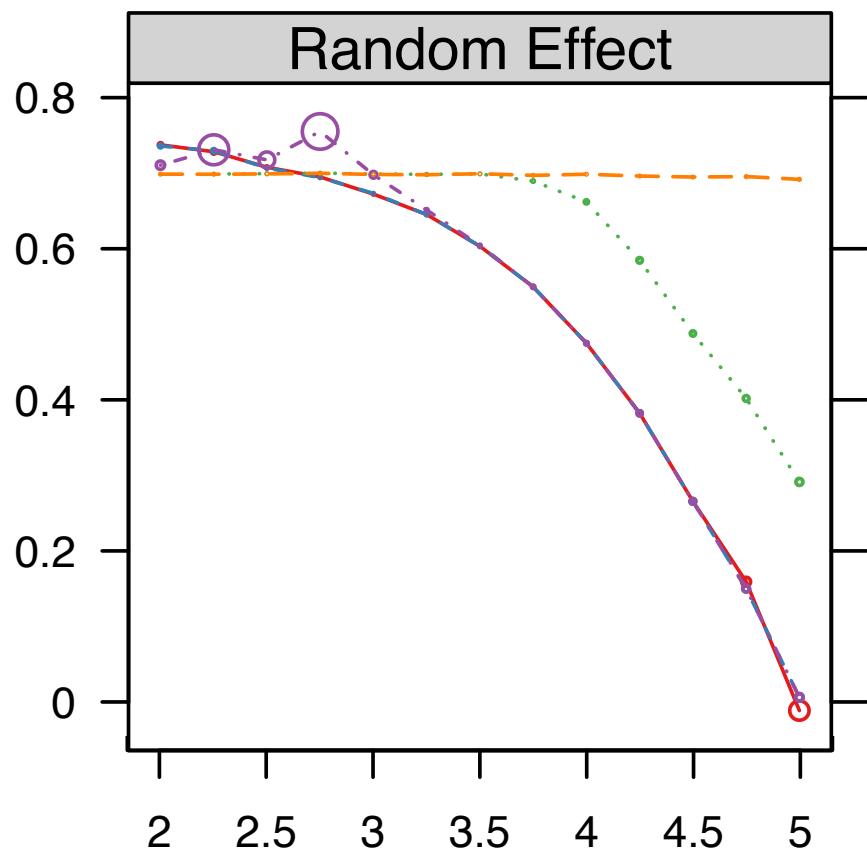
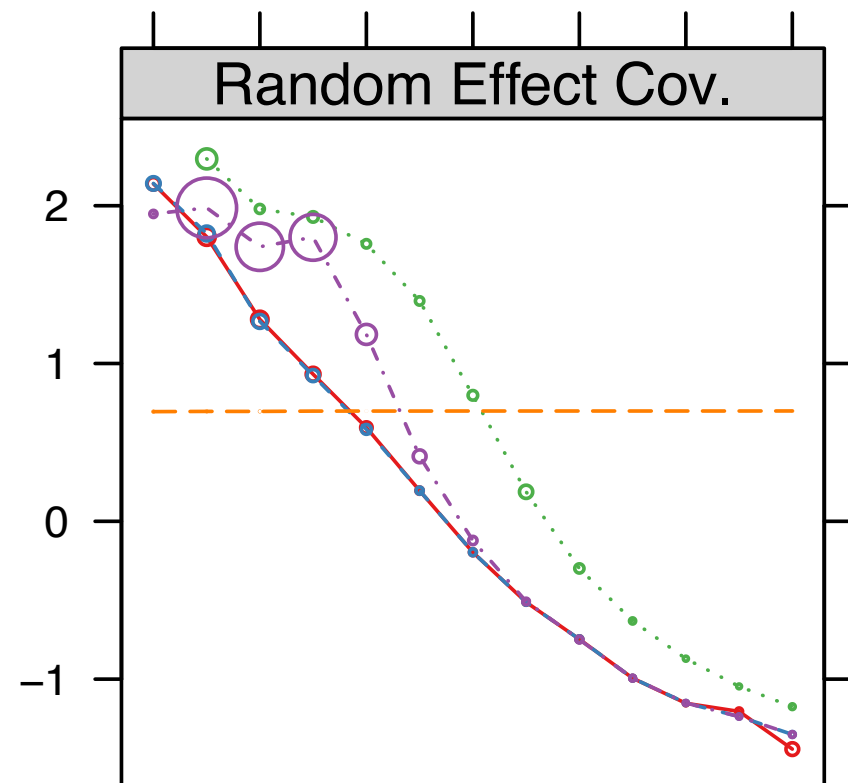
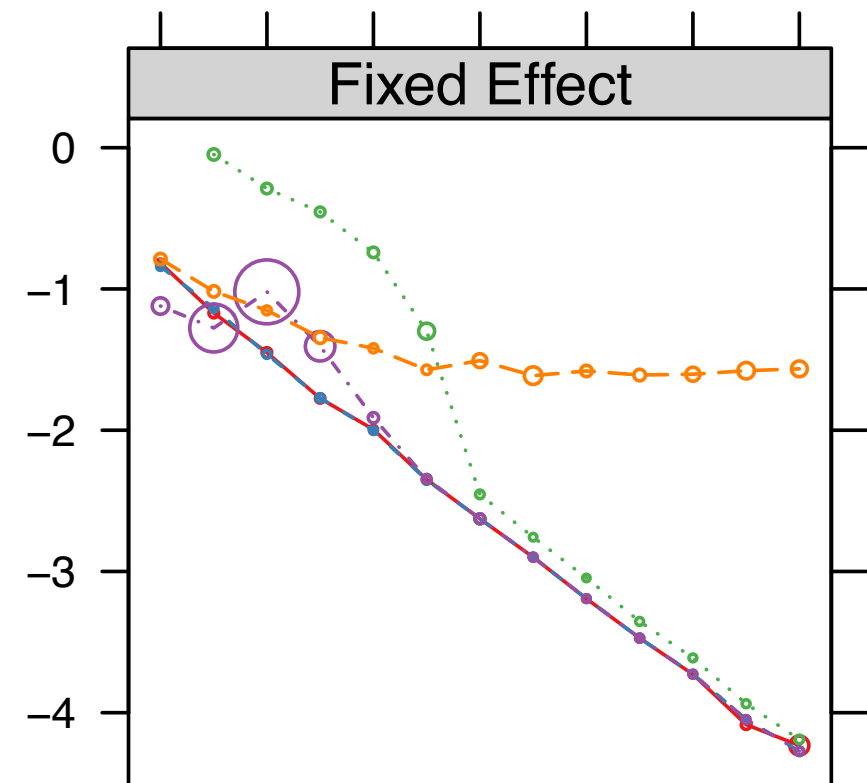
# Linear Model: Time



# Linear Model: Accuracy



Log<sub>10</sub> Mean Loss




Method



Log<sub>10</sub> Samples

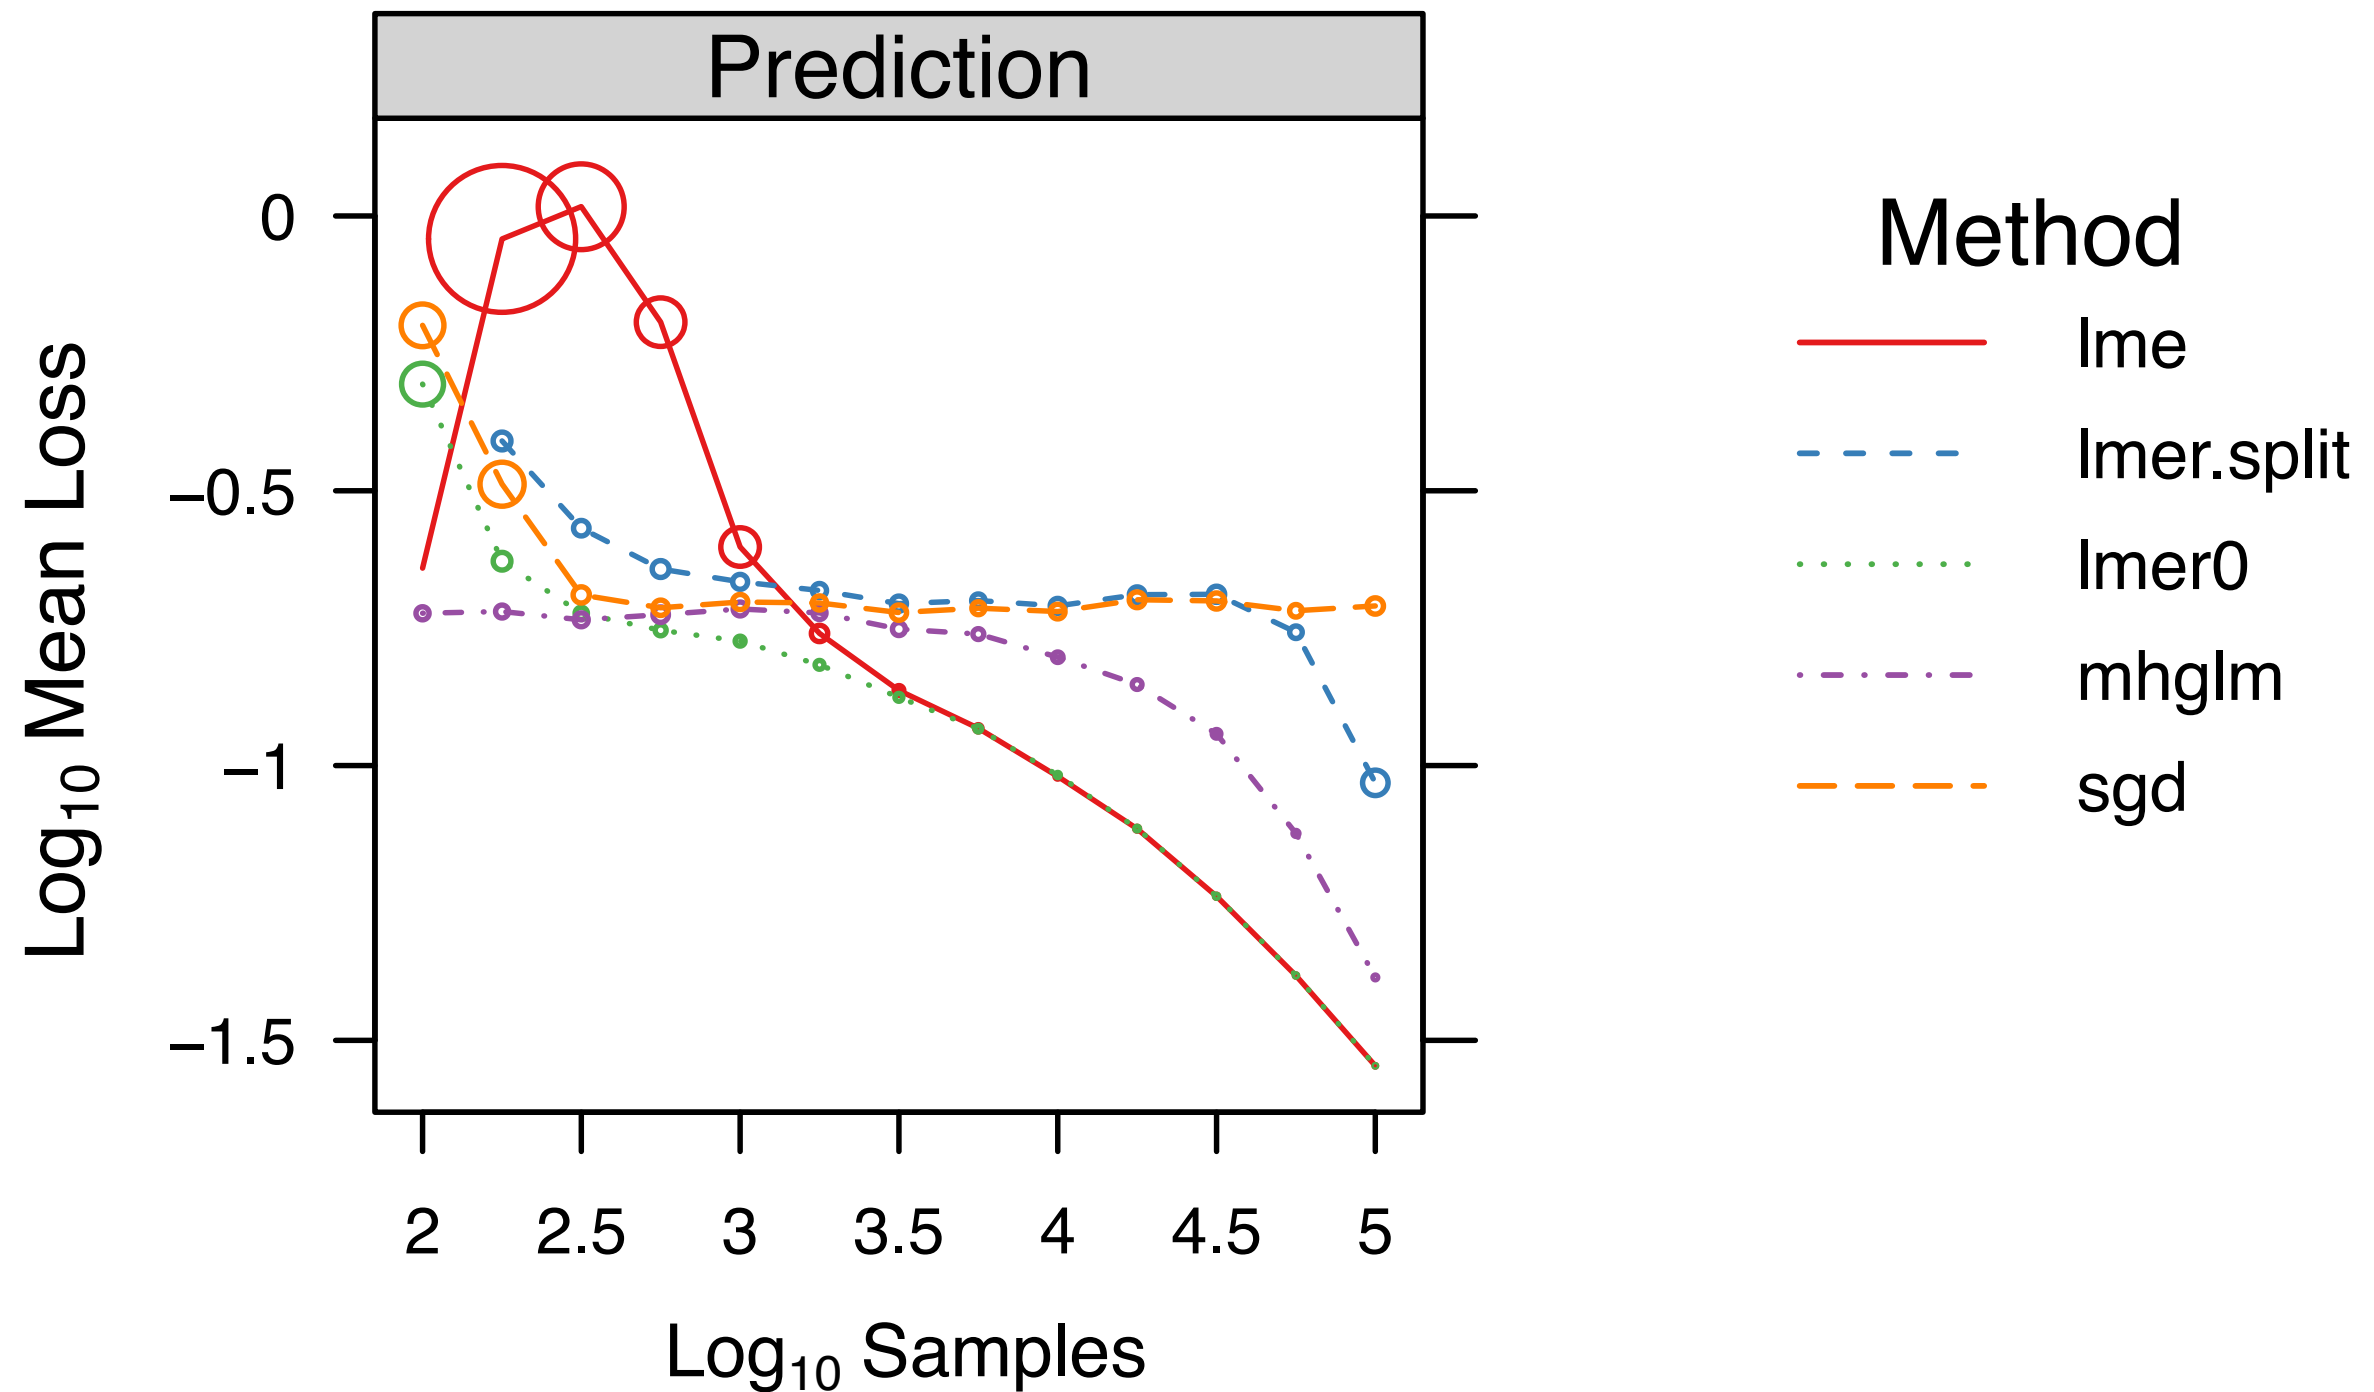
# Performance in Hierarchical Logistic Model Simulations

# Logistic Model: Time










# Logistic Model: Accuracy



Log<sub>10</sub> Mean Loss



Method



Log<sub>10</sub> Samples

# Performance in Practice

# Application: MovieLens 10M

MovieLens 10M dataset:

- 10 million ratings
- 70 thousand users
- 10 thousand movies
- Predictors: genre, item popularity, user mood

Time to fit with *glmer*:  
10 hours

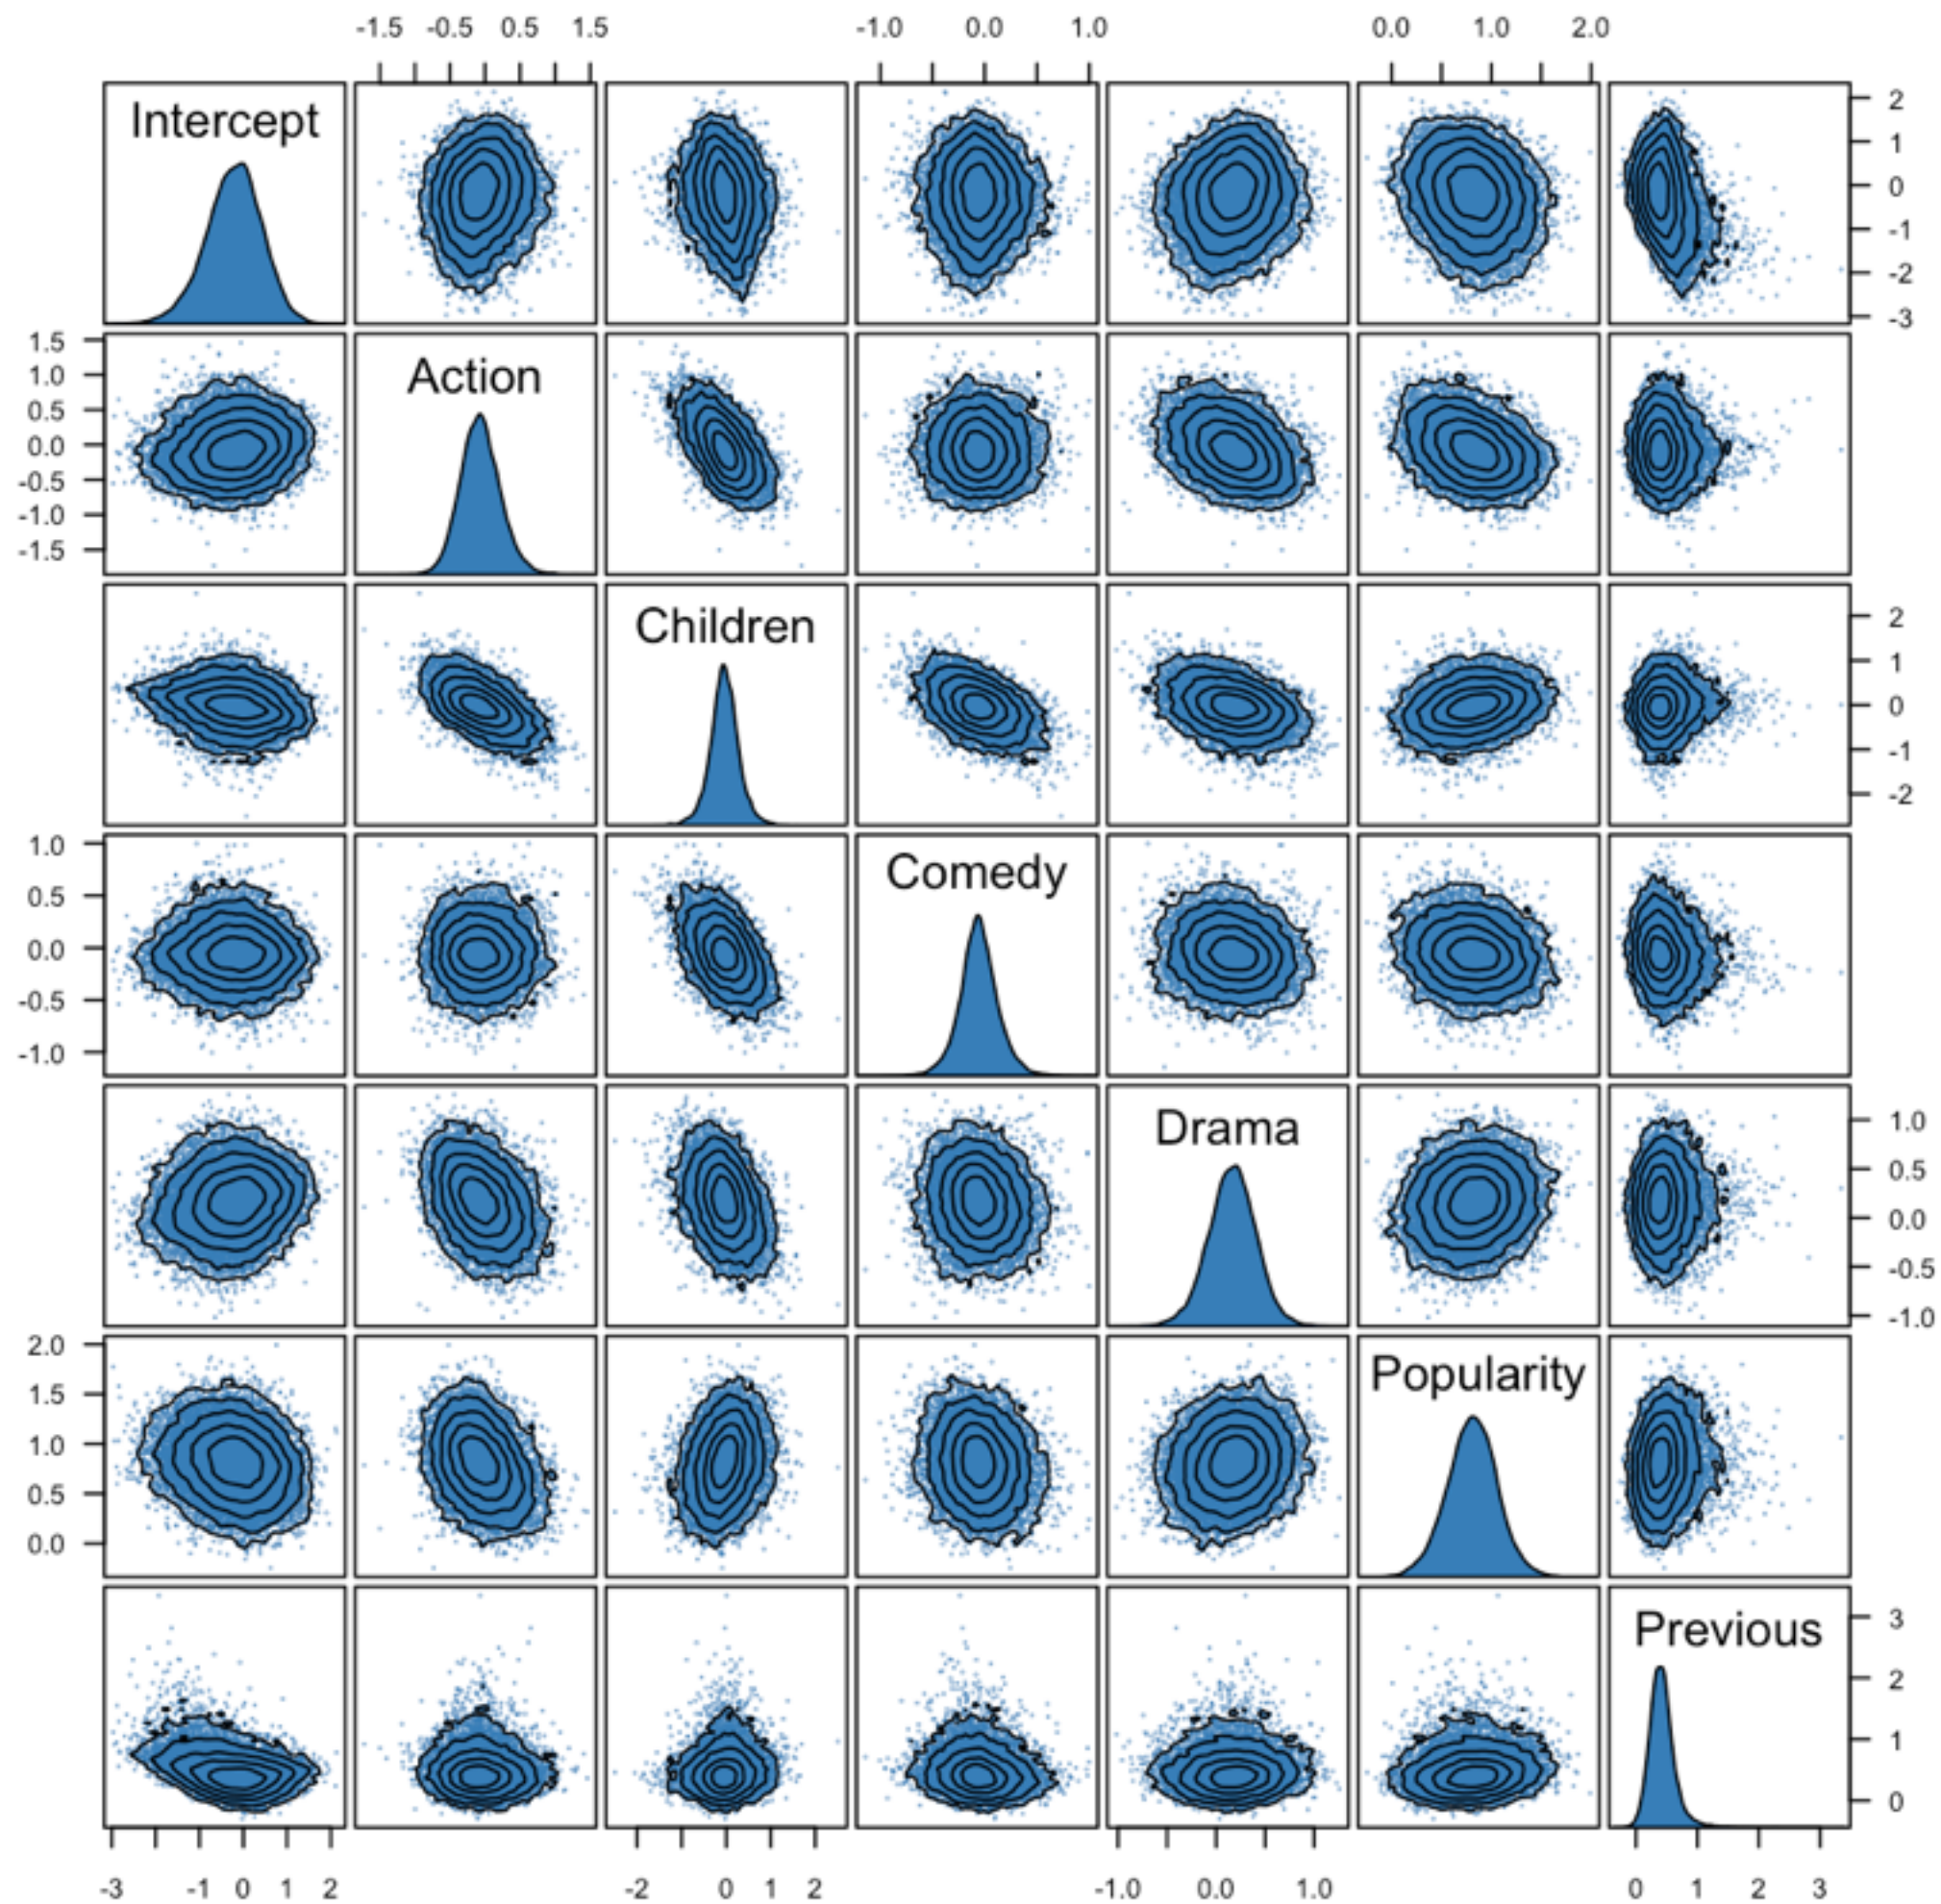
Time to fit with *mhglm*:  
10 minutes

# Response

- Given that user  $i$  rates movie  $j$ :  
 $y_{ij}$  = rating is positive (4 or 5 stars)
- Ratings per user: 140 (ranges from 10 to 1000)

# Predictors

- **Genre:** Action, Children, Comedy, or Drama
- **Popularity:**  $\text{logit}(\text{current popularity of movie})$ , computed from 30 most recent ratings
- **Previous:** indicator of whether user's previous review was positive




# Out-of-Sample Performance

Table 3: MovieLens test set error.

Method	Loss Function		
	Log	Squared Error	Misclassification
Hierarchical	0.55	0.18	0.28
Local	0.57	0.19	0.29
Global	0.59	0.20	0.32





# Summary: Moment-Based Estimation

- Fast
- Theoretically Sound
- Works In Practice

# R package: mbest

```
mhglm(y ~ genre + popularity  
      + (1 + genre | user),  
      family=binomial)
```

Available on CRAN

**Thank you!**